

week3R

Emmanuel Titi

2022-06-11

Loading our data.

```
df=read.csv("http://bit.ly/CarreFourDataset")
summary(df)
```

```
## Invoice.ID      Branch      Customer.type      Gender
## Length:1000    Length:1000    Length:1000    Length:1000
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
##
##
##
## Product.line    Unit.price    Quantity    Tax
## Length:1000    Min.      :10.08    Min.       : 1.00    Min.       : 0.5085
## Class :character 1st Qu.:32.88    1st Qu.: 3.00    1st Qu.: 5.9249
## Mode  :character Median :55.23    Median : 5.00    Median :12.0880
##                  Mean  :55.67    Mean  : 5.51    Mean  :15.3794
##                  3rd Qu.:77.94    3rd Qu.: 8.00    3rd Qu.:22.4453
##                  Max.   :99.96    Max.   :10.00    Max.   :49.6500
##
## Date            Time            Payment      cogs
## Length:1000    Length:1000    Length:1000    Min.      : 10.17
## Class :character Class :character Class :character 1st Qu.:118.50
## Mode  :character Mode  :character Mode  :character Median :241.76
##                  Mean  :307.59
##                  3rd Qu.:448.90
##                  Max.   :993.00
##
## gross.margin.percentage gross.income    Rating    Total
## Min.      :4.762    Min.      : 0.5085    Min.      : 4.000    Min.      : 10.68
## 1st Qu.:4.762    1st Qu.: 5.9249    1st Qu.: 5.500    1st Qu.: 124.42
## Median :4.762    Median :12.0880    Median : 7.000    Median : 253.85
## Mean  :4.762    Mean  :15.3794    Mean  : 6.973    Mean  : 322.97
## 3rd Qu.:4.762    3rd Qu.:22.4453    3rd Qu.: 8.500    3rd Qu.: 471.35
## Max.   :4.762    Max.   :49.6500    Max.   :10.000    Max.   :1042.65
```

Removing features not required

```
df1<-dplyr::select(df,-c("Invoice.ID","Date","Time"))
```

Onehot-encoding out categorical variables.

```
#dummify the data
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
df_dummy<-dummyVars("~.",data=df1)
df2<-data.frame(predict(df_dummy,newdat=df1))
head(df2)
```

```
##   BranchA BranchB BranchC Customer.typeMember Customer.typeNormal GenderFemale
## 1      1      0      0      1      0      1
## 2      0      0      1      0      1      1
## 3      1      0      0      0      1      0
## 4      1      0      0      1      0      0
## 5      1      0      0      0      1      0
## 6      0      0      1      0      1      0
##   GenderMale Product.lineElectronic.accessories Product.lineFashion.accessories
## 1          0                                0                                0
## 2          0                                1                                0
## 3          1                                0                                0
## 4          1                                0                                0
## 5          1                                0                                0
## 6          1                                1                                0
##   Product.lineFood.and.beverages Product.lineHealth.and.beauty
## 1                               0                               1
## 2                               0                               0
## 3                               0                               0
## 4                               0                               1
## 5                               0                               0
## 6                               0                               0
##   Product.lineHome.and.lifestyle Product.lineSports.and.travel Unit.price
## 1                               0                               0      74.69
## 2                               0                               0      15.28
## 3                               1                               0      46.33
## 4                               0                               0      58.22
## 5                               0                               1      86.31
## 6                               0                               0      85.39
##   Quantity      Tax PaymentCash PaymentCredit.card PaymentEwallet   cogs
## 1       7 26.1415          0          0          1 522.83
## 2       5  3.8200          1          0          0  76.40
## 3       7 16.2155          0          1          0 324.31
## 4       8 23.2880          0          0          1 465.76
## 5       7 30.2085          0          0          1 604.17
## 6       7 29.8865          0          0          1 597.73
##   gross.margin.percentage gross.income Rating   Total
## 1          4.761905      26.1415    9.1 548.9715
## 2          4.761905       3.8200    9.6  80.2200
## 3          4.761905      16.2155    7.4 340.5255
## 4          4.761905      23.2880    8.4 489.0480
## 5          4.761905      30.2085    5.3 634.3785
## 6          4.761905      29.8865    4.1 627.6165
```

We need to separate the encoded data and the original numerical data so as to scale without getting an error

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
df_scale<-df2%>%select( Total,Rating, gross.income,,cogs
                        ,Quantity,Tax,Unit.price)
df_encod<-df2%>%select(-Total,-Rating, -gross.income,
                      -gross.margin.percentage,-cogs
                      ,-Quantity,-Tax,-Unit.price)
```

scaling our data

```
library(tibble)
scaled<-scale(as.data.frame(df_scale,center = TRUE))
enframe(scaled)
```

```
## # A tibble: 1,000 x 2
##   name value[, "Total"] [, "Rating"] [, "gross.income"] [, "cogs"] [, "Quantity"]
##   <chr>          <dbl>          <dbl>          <dbl>          <dbl>          <dbl>
## 1 1              0.919            1.24            0.919            0.919            0.510
## 2 2             -0.987            1.53           -0.987           -0.987           -0.174
## 3 3              0.0714           0.249           0.0714           0.0714            0.510
## 4 4              0.675            0.831           0.675            0.675            0.852
## 5 5              1.27           -0.973           1.27             1.27            0.510
## 6 6              1.24           -1.67            1.24             1.24            0.510
## 7 7              0.450           -0.682           0.450            0.450            0.168
## 8 8              1.83            0.598            1.83             1.83            1.54
## 9 9             -1.00            0.132           -1.00            -1.00           -1.20
## 10 10            -0.611           -0.624           -0.611           -0.611           -0.859
## # ... with 990 more rows, and 1 more variable: value[6:7] <dbl>
```

```
head(scaled)
```

```
##           Total      Rating gross.income      cogs  Quantity      Tax
## 1  0.91914693  1.2378240  0.91914693  0.91914693  0.5096752  0.91914693
## 2 -0.98723557  1.5287619 -0.98723557 -0.98723557 -0.1744526 -0.98723557
## 3  0.07141032  0.2486355  0.07141032  0.07141032  0.5096752  0.07141032
## 4  0.67544187  0.8305111  0.67544187  0.67544187  0.8517391  0.67544187
## 5  1.26649176 -0.9733034  1.26649176  1.26649176  0.5096752  1.26649176
```

```
## 6 1.23899114 -1.6715541 1.23899114 1.23899114 0.5096752 1.23899114
## Unit.price
## 1 0.71780097
## 2 -1.52454035
## 3 -0.35260468
## 4 0.09616553
## 5 1.15638044
## 6 1.12165642
```

Lets now join the two the dataframes

```
#concatinate the scaled and dfencode
newdf<-cbind(scaled,df_encod)
head(newdf)
```

```
##      Total      Rating gross.income      cogs  Quantity      Tax
## 1  0.91914693  1.2378240  0.91914693  0.91914693  0.5096752  0.91914693
## 2 -0.98723557  1.5287619 -0.98723557 -0.98723557 -0.1744526 -0.98723557
## 3  0.07141032  0.2486355  0.07141032  0.07141032  0.5096752  0.07141032
## 4  0.67544187  0.8305111  0.67544187  0.67544187  0.8517391  0.67544187
## 5  1.26649176 -0.9733034  1.26649176  1.26649176  0.5096752  1.26649176
## 6  1.23899114 -1.6715541  1.23899114  1.23899114  0.5096752  1.23899114
##      Unit.price BranchA BranchB BranchC Customer.typeMember Customer.typeNormal
## 1  0.71780097      1      0      0      1      0
## 2 -1.52454035      0      0      1      0      1
## 3 -0.35260468      1      0      0      0      1
## 4  0.09616553      1      0      0      1      0
## 5  1.15638044      1      0      0      0      1
## 6  1.12165642      0      0      1      0      1
##      GenderFemale GenderMale Product.lineElectronic.accessories
## 1      1      0      0
## 2      1      0      1
## 3      0      1      0
## 4      0      1      0
## 5      0      1      0
## 6      0      1      1
##      Product.lineFashion.accessories Product.lineFood.and.beverages
## 1      0      0
## 2      0      0
## 3      0      0
## 4      0      0
## 5      0      0
## 6      0      0
##      Product.lineHealth.and.beauty Product.lineHome.and.lifestyle
## 1      1      0
## 2      0      0
## 3      0      1
## 4      1      0
## 5      0      0
## 6      0      0
##      Product.lineSports.and.travel PaymentCash PaymentCredit.card PaymentEwallet
## 1      0      0      0      1
## 2      0      1      0      0
## 3      0      0      1      0
```

```
## 4          0          0          0          1
## 5          1          0          0          1
## 6          0          0          0          1
```

Lets run our pca function

```
#PCA
dfPCA<-prcomp(newdf)
#head(dfPCA)
```

Lets make a short descriptive analysis of the PCs


```
summary(dfPCA)
```

```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  2.2191 1.00198 0.99639 0.72621 0.69392 0.59438 0.58110
## Proportion of Variance 0.4843 0.09873 0.09763 0.05186 0.04735 0.03474 0.03321
## Cumulative Proportion 0.4843 0.58300 0.68064 0.73250 0.77985 0.81460 0.84780
##          PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation  0.56903 0.55868 0.41794 0.41052 0.40776 0.40066 0.3905
## Proportion of Variance 0.03184 0.03069 0.01718 0.01657 0.01635 0.01579 0.0150
## Cumulative Proportion 0.87965 0.91034 0.92752 0.94409 0.96044 0.97623 0.9912
##          PC15     PC16     PC17     PC18     PC19     PC20
## Standard deviation  0.29863 4.594e-16 4.12e-16 2.817e-16 2.471e-16 1.986e-16
## Proportion of Variance 0.00877 0.000e+00 0.00e+00 0.000e+00 0.000e+00 0.000e+00
## Cumulative Proportion 1.00000 1.000e+00 1.00e+00 1.000e+00 1.000e+00 1.000e+00
##          PC21     PC22     PC23
## Standard deviation  1.683e-16 1.001e-16 3.444e-17
## Proportion of Variance 0.000e+00 0.000e+00 0.000e+00
## Cumulative Proportion 1.000e+00 1.000e+00 1.000e+00
```

We can start to visualize our PCs using the plot function

```
plot(dfPCA,type="l")
```


dfPCA



```
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
fviz_eig(dfPCA ,addlabels = TRUE,ylim=c(0,50))
```




The scree plot shows percentage of variance contributed by each principal component. We might want to stop at 6 since already 81% of variance is represented.

```
df3<-cbind(df,(dfPCA$x[,1:6]))
```

Here we have now combined our original data set to the 1st 6 components

```
library(ggplot2)
ggplot(df3,aes(PC1,PC3,col=Branch,fill=Branch))+
  stat_ellipse(geom="polygon",col="black",alpha=0.5)+
  geom_point(shape=21,col="black")
```



The plot shows the relationship Between PC1 and 2 separated/highlighted branch-wise.

Graph of variables

Here we use the `get_pca_var()` function that provides a list of matrices containing all results for active variables.

```
var_df<-get_pca_var(dfPCA)
var_df
```

```
## Principal Component Analysis Results for variables
## =====
##   Name      Description
## 1 "$coord"   "Coordinates for the variables"
## 2 "$cor"     "Correlations between variables and dimensions"
## 3 "$cos2"    "Cos2 for the variables"
## 4 "$contrib" "contributions of the variables"
```

We can now visualize variables and make conclusions

Correlation circle

The correlation between variables and principal component is used as the coordinates of the variable on the PC

```
#coordinates of variables
head(var_df$coord,6)
```


	Dim.1	Dim.2	Dim.3	Dim.4	Dim.5
## Total	-0.99782094	0.004441866	-0.001602519	0.006105305	-0.0009938612
## Rating	0.04147205	0.903838165	-0.424135766	-0.010112460	-0.0114491551
## gross.income	-0.99782094	0.004441866	-0.001602519	0.006105305	-0.0009938612
## cogs	-0.99782094	0.004441866	-0.001602519	0.006105305	-0.0009938612
## Quantity	-0.72024050	-0.267716439	-0.605532437	0.031880602	-0.0339532773
## Tax	-0.99782094	0.004441866	-0.001602519	0.006105305	-0.0009938612
	Dim.6	Dim.7	Dim.8	Dim.9	Dim.10
## Total	0.001180199	-0.002108020	0.001165783	-0.005950078	-0.001750018
## Rating	0.005177509	-0.023970262	0.004548973	0.022249491	-0.001087060
## gross.income	0.001180199	-0.002108020	0.001165783	-0.005950078	-0.001750018
## cogs	0.001180199	-0.002108020	0.001165783	-0.005950078	-0.001750018
## Quantity	-0.008022062	0.005916874	-0.003612045	0.024485248	-0.005213397
## Tax	0.001180199	-0.002108020	0.001165783	-0.005950078	-0.001750018
	Dim.11	Dim.12	Dim.13	Dim.14	Dim.15
## Total	-0.001496901	-0.0001471787	0.002459825	-0.001946532	-0.065077266
## Rating	0.006446994	0.0017339896	0.005330317	-0.003339676	-0.005451258
## gross.income	-0.001496901	-0.0001471787	0.002459825	-0.001946532	-0.065077266
## cogs	-0.001496901	-0.0001471787	0.002459825	-0.001946532	-0.065077266
## Quantity	-0.006063208	0.0005932810	-0.014573120	0.004593227	0.199331606
## Tax	-0.001496901	-0.0001471787	0.002459825	-0.001946532	-0.065077266
	Dim.16	Dim.17	Dim.18	Dim.19	
## Total	1.151192e-17	-1.472107e-16	-2.859765e-17	-1.219424e-17	
## Rating	-9.563348e-33	5.717951e-32	0.000000e+00	2.571745e-33	
## gross.income	-2.828127e-17	2.637773e-16	3.739402e-17	-3.174842e-17	
## cogs	1.548519e-17	-2.294141e-16	-2.631499e-17	2.708411e-17	
## Quantity	-6.534954e-32	1.029231e-31	2.443142e-32	-8.572485e-34	
## Tax	1.284155e-18	1.128475e-16	1.751861e-17	1.685855e-17	
	Dim.20	Dim.21	Dim.22	Dim.23	
## Total	8.692806e-18	-8.429379e-18	7.634883e-17	-5.211008e-18	
## Rating	3.444621e-33	-4.670443e-33	1.667131e-32	4.780040e-33	
## gross.income	2.750408e-20	-9.840235e-18	-1.260040e-17	-1.832186e-17	
## cogs	6.404734e-18	2.768578e-17	-6.120422e-17	-4.383107e-18	
## Quantity	-8.267092e-33	-2.130890e-32	8.335654e-33	-1.027709e-32	
## Tax	-1.512504e-17	-9.416165e-18	-2.544203e-18	2.791598e-17	

Results above shows how the variables correlate to the PCs.

Quality of representation

Quality of representation on a factor map is **cos2** (squared cosine ,squared coordinates)

```
head(var_df$cos2,6)
```

	Dim.1	Dim.2	Dim.3	Dim.4	Dim.5
## Total	0.995646625	1.973017e-05	2.568066e-06	3.727475e-05	9.877601e-07
## Rating	0.001719931	8.169234e-01	1.798911e-01	1.022619e-04	1.310832e-04
## gross.income	0.995646625	1.973017e-05	2.568066e-06	3.727475e-05	9.877601e-07
## cogs	0.995646625	1.973017e-05	2.568066e-06	3.727475e-05	9.877601e-07
## Quantity	0.518746385	7.167209e-02	3.666695e-01	1.016373e-03	1.152825e-03
## Tax	0.995646625	1.973017e-05	2.568066e-06	3.727475e-05	9.877601e-07
	Dim.6	Dim.7	Dim.8	Dim.9	Dim.10
## Total	1.392869e-06	4.443749e-06	1.359051e-06	3.540343e-05	3.062563e-06
## Rating	2.680660e-05	5.745735e-04	2.069315e-05	4.950398e-04	1.181699e-06


```
## gross.income 1.392869e-06 4.443749e-06 1.359051e-06 3.540343e-05 3.062563e-06
## cogs         1.392869e-06 4.443749e-06 1.359051e-06 3.540343e-05 3.062563e-06
## Quantity    6.435348e-05 3.500940e-05 1.304687e-05 5.995274e-04 2.717950e-05
## Tax         1.392869e-06 4.443749e-06 1.359051e-06 3.540343e-05 3.062563e-06
##              Dim.11      Dim.12      Dim.13      Dim.14      Dim.15
## Total       2.240713e-06 2.166157e-08 6.050738e-06 3.788987e-06 4.235051e-03
## Rating      4.156373e-05 3.006720e-06 2.841228e-05 1.115344e-05 2.971621e-05
## gross.income 2.240713e-06 2.166157e-08 6.050738e-06 3.788987e-06 4.235051e-03
## cogs        2.240713e-06 2.166157e-08 6.050738e-06 3.788987e-06 4.235051e-03
## Quantity    3.676250e-05 3.519823e-07 2.123758e-04 2.109773e-05 3.973309e-02
## Tax         2.240713e-06 2.166157e-08 6.050738e-06 3.788987e-06 4.235051e-03
##              Dim.16      Dim.17      Dim.18      Dim.19      Dim.20
## Total       1.325244e-34 2.167099e-32 8.178254e-34 1.486995e-34 7.556488e-35
## Rating      9.145762e-65 3.269496e-63 0.000000e+00 6.613874e-66 1.186542e-65
## gross.income 7.998303e-34 6.957846e-32 1.398313e-33 1.007962e-33 7.564745e-40
## cogs        2.397912e-34 5.263082e-32 6.924785e-34 7.335488e-34 4.102062e-35
## Quantity    4.270563e-63 1.059317e-62 5.968942e-64 7.348749e-67 6.834480e-65
## Tax         1.649054e-36 1.273455e-32 3.069018e-34 2.842109e-34 2.287670e-34
##              Dim.21      Dim.22      Dim.23
## Total       7.105442e-35 5.829143e-33 2.715461e-35
## Rating      2.181304e-65 2.779325e-64 2.284879e-65
## gross.income 9.683022e-35 1.587701e-34 3.356906e-34
## cogs        7.665023e-34 3.745957e-33 1.921163e-35
## Quantity    4.540691e-64 6.948312e-65 1.056185e-64
## Tax         8.866416e-35 6.472970e-36 7.793018e-34
```

lets visualize our results on cos2


```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
corrplot(var_df$cos2,is.corr=FALSE)
```



```
fviz_cos2(dfPCA,choice="var",xes=1:2)
```



A high cos2 indicates good representation of the variable on the principal component.

We can hence conclude that the variables below are the more important features in our data.

```
str(newdf[,1:7])
```

```
## 'data.frame': 1000 obs. of 7 variables:
## $ Total : num 0.9191 -0.9872 0.0714 0.6754 1.2665 ...
## $ Rating : num 1.238 1.529 0.249 0.831 -0.973 ...
## $ gross.income: num 0.9191 -0.9872 0.0714 0.6754 1.2665 ...
## $ cogs : num 0.9191 -0.9872 0.0714 0.6754 1.2665 ...
## $ Quantity : num 0.51 -0.174 0.51 0.852 0.51 ...
## $ Tax : num 0.9191 -0.9872 0.0714 0.6754 1.2665 ...
## $ Unit.price : num 0.7178 -1.5245 -0.3526 0.0962 1.1564 ...
```

FEATURE SELECTION

Random forest method

```
#libraries
```

```
library(party)
```

```
## Loading required package: grid
```

```
## Loading required package: mvtnorm
```

```

## Loading required package: modeltools

## Loading required package: stats4

## Loading required package: strucchange

## Loading required package: zoo

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric

## Loading required package: sandwich

library(randomForest)

## randomForest 4.7-1.1

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:dplyr':
##
##      combine

## The following object is masked from 'package:ggplot2':
##
##      margin

cf1 <- cforest(Total ~ . , data= newdf, control=cforest_unbiased(mtry=2,ntree=50))

#feature selection
varimp(cf1)

##
##           Rating           gross.income
##      3.384917e-04      3.907661e-01
##           cogs           Quantity
##      3.301923e-01      1.319855e-01
##           Tax           Unit.price
##      2.620213e-01      9.444527e-02
##           BranchA           BranchB
##      -1.082237e-03      7.928378e-04
##           BranchC      Customer.typeMember
##      -7.280935e-04      -1.579422e-03
##      Customer.typeNormal      GenderFemale

```

```
##          -2.931581e-03          4.768697e-04
##          GenderMale Product.lineElectronic.accessories
##          3.305157e-04          -3.461942e-04
## Product.lineFashion.accessories Product.lineFood.and.beverages
##          -6.611270e-04          -4.664718e-03
## Product.lineHealth.and.beauty Product.lineHome.and.lifestyle
##          -1.164853e-03          -1.494883e-03
## Product.lineSports.and.travel PaymentCash
##          5.854174e-05          -1.480815e-03
##          PaymentCredit.card PaymentEwallet
##          7.063618e-04          5.464844e-04
```

```
# get variable importance, based on mean decrease in accuracy
```